

Determinants of Advertised Pay in the US Energy and Data Center Sector

Evidence from employer-published job postings

Alexander J. Henderson

Kelley School of Business, Indiana University

543 postings · 100 employers · collected 20–30 September 2026

AN OBSERVATION IN THE SAMPLE

Eversource Energy · the employer's own careers site

Associate Analyst, Transmission (Hybrid Schedule)

Apply

Westwood, MA
Hooksett, NH
Berlin, CT

Full time
Posted 19 Days Ago
R-031417

Pay range, stated further down the posting

\$71,610 – \$79,570

Lists MA, NH and CT: covered through Massachusetts (law in force since October 29, 2025).

Midpoint used in the pay model: \$75,590

Screenshot of the employer's posting (October 2026); pay range and coverage from the collected record.

What in a job posting predicts the pay an employer advertises?

1 Pay transparency is spreading

13 US jurisdictions now require a pay range in the posting itself; Colorado was first, in 2021.

2 Energy hiring is changing

Data-center construction is pulling analytical talent into utilities, developers and grid operators.

3 Students can benchmark offers

For a graduating student, the posted range is the first price signal of a career.

The author as the subject

Industries and roles were benchmarked to the author's résumé: the sectors with the most experience and interest, and roles the author holds or could hold. The same design extends to any student or industry.

RELATED WORK


Arnold, Quach & Taska (2025, NBER WP 34480): Colorado's posting law raised the share of postings stating pay by about 30 points.

Cullen & Pakzad-Hurson (2023, Econometrica): equilibrium effects of transparency on wages.

Cullen (2024, Journal of Economic Perspectives): is pay transparency good? A review.

This study: one sector, all US states, 729 postings; disclosure is modelled as an outcome, not assumed.

13 jurisdictions required a pay range in the posting during the collection window



■ Posting mandate in force (year it took effect)


▤ Takes effect after the window (CT, 2026-10-01)

- 1 Coded by date**
Each posting is coded against the laws in force on the day it was seen; CT's rule starts 2026-10-01, after the window.
- 2 Coverage follows the job**
A posting is covered if any listed work location is in a mandate state (14% list more than one).
- 3 Elsewhere, disclosure is a choice**
Without a law, stating pay is voluntary, so postings that state pay are a selected group.


The question for the data: does disclosure differ where the law requires it?

The sample: analytical roles in energy, utilities and data centers

In-scope postings by industry (N = 729)



In-scope postings by role family (N = 729)



Benchmarked to the author's résumé. Industries where the author has the most experience and interest; roles the author is, or could become, qualified for. Every seniority level is kept, because seniority is a regressor; internships are excluded (a different contract and pay regime).

Postings come straight from the job boards employers publish to

Source	Employers with postings	Postings retrieved*	In scope (unique)	State pay
Greenhouse	67	5,605	157	125
Workday	56	4,591	389	286
Ashby	25	4,010	111	82
Lever	21	2,870	35	27
USAJOBS (federal)	6	220	3	3
Workable	8	130	20	14
SmartRecruiters	7	68	14	6
Recruitee	1	1	0	0
Total	191	17,495	729	543

*Every sighting on each of the 7 collection days (2026-09-20 to 2026-09-30). USAJOBS is the federal government's job API. Employer counts sum across sources.

407

employers searched in the sampling frame

185 had a readable public board on the final run; the rest use systems with no open feed, or none at all.

NOT USED: TERMS PROHIBIT AUTOMATED COLLECTION

LinkedIn

Indeed

Handshake


iCIMS

Oracle Cloud HCM

UKG


careers.electric.coop

From 17,495 retrieved postings to an estimation sample of 543



A posting can fail more than one screen; each is removed at the first it fails. Duplicates are matched on employer, title and location.

100 employers in 31 states; the largest supplies 9.0% of the sample




Postings that state pay: 1–5 6–20 21–60 61+

Plus 45 nationwide-remote postings with no single state.



Employers stating pay, by industry



Two models: what predicts pay, and what predicts stating it

Pay model · N = 543 postings that state pay

$$\ln(\text{Pay}_{ij}) = \alpha + \beta_1 \text{Seniority}_{ij} + \beta_2 \text{Experience}_{ij} + X'_{ij}\gamma + \varepsilon_{ij}$$

Pay is the midpoint of the advertised range, in annual US dollars. X: degree, skills, region, mandate, industry, role family (15 regressors in all).

Disclosure model · N = 729 in-scope postings

$$\Pr(\text{Pay stated}_{ij}) = \delta + \theta \text{Mandate}_{ij} + Z'_{ij}\lambda + u_{ij}$$

A linear probability model. Z controls for seniority, remote eligibility, data-center operator and census region.

1

Pre-registered

Seven hypotheses and their signs were committed before the national sample was collected.

2

Specification by sample size

36.2 observations per regressor (rule: at least 20), so the extended model (27 regressors) is estimated too.

3

Few-cluster inference

Standard errors clustered by employer (100); every p-value read from a wild cluster bootstrap (Rademacher, 9,999 draws).


4

Four robustness checks

Price-adjusted pay; without nationwide-remote postings; without Crusoe; without the federal rows.


i indexes postings, j employers. Associational design: one cross-section with no time variation, so no difference-in-differences.

Pay is stated in 91% of postings where the law requires it, 48% elsewhere



43.8 percentage-point gap
share stating pay, mandate minus no mandate


THE GAP UNDER EVERY CUT OF THE SAMPLE




Associational, not causal: a single cross-section; employers in mandate states differ in other ways.

With controls, a mandate is still associated with a 39-point higher chance of stating pay

Change in the probability of stating pay, percentage points with 95% CI (N = 729)




Share stating pay, by region



The South has the lowest disclosure (42%) and 32% of its postings covered; with region held fixed the gap goes from 43.8 to 38.8 points.

Standard errors clustered by employer. $R^2 = 0.335$. Crimson: $p < 0.05$ (clustered; this model has no bootstrap). Region effects are relative to the Midwest.

Seniority, required experience and a stated cloud skill predict advertised pay under every check



HOW TO READ IT

- **Holds under every check**
seniority, required experience and a stated cloud skill
- **Tentative or nominal only**
passes the bootstrap, fails a check
- **Overtaken by the bootstrap**
clustered $p < 0.05$, bootstrap $p \geq 0.05$
- **Not significant**

Core model, $N = 543$, 100 employer clusters, $R^2 = 0.589$. Dot: % effect, $e^{\beta} - 1$. Bar: 95% cluster-robust CI.

In dollars: one seniority step is worth about \$17,000 a year

+11.2%

per seniority step

≈ \$17,000 a year at mean pay

Pre-registered

+2.5%

per year of required experience

≈ \$3,900 a year at mean pay


Pre-registered

+12.5%

for a stated cloud skill

≈ \$19,000 a year at mean pay

Exploratory



Mean advertised pay \$151,839 (median \$144,360). Predictions are geometric means from the core model: $\exp(\text{mean } \ln \text{ pay} + \beta \times (\text{rank} - \text{mean rank}))$.

Only three results hold under the bootstrap and all four checks

Attribute	Clustered SE	Wild bootstrap	Without nationwide-remote	Price-adjusted pay	Without Crusoe	Without federal rows	Verdict
Seniority (per step)	<0.001	<0.001	<0.001	<0.001	<0.001	<0.001	Holds everywhere
Required experience (per year)	<0.001	<0.001	<0.001	<0.001	<0.001	<0.001	Holds everywhere
Cloud skill (AWS, Azure, GCP)	<0.001	<0.001	0.003	<0.001	0.002	0.001	Holds everywhere
West (vs. Midwest)	0.001	0.011	0.022	0.149	0.012	0.015	Nominal only
Northeast (vs. Midwest)	0.005	0.019	0.023	0.706	0.031	0.018	Nominal only
States an experience minimum	0.020	0.023	0.039	0.017	0.033	0.030	Tentative
STEM degree	0.031	0.043	0.028	0.036	0.086	0.051	Tentative
Requires a degree	0.055	0.074	0.104	0.114	0.156	0.078	Not significant
ML/AI skill	0.033	0.074	0.011	0.002	0.207	0.069	Overtaken
AI/ML role family	0.048	0.083	0.070	0.019	0.087	0.080	Overtaken
South (vs. Midwest)	0.063	0.154	0.189	0.306	0.072	0.141	Not significant
Data-center operator	0.160	0.584	0.677	0.212	0.365	0.589	Not significant
Advertised hourly	0.535	0.644	0.684	0.088	0.623	0.661	Not significant
Pay-transparency mandate	0.730	0.744	0.923	0.086	0.811	0.692	Not significant
Remote-eligible	0.772	0.787	0.675	0.844	0.507	0.719	Not significant


Cells: p-values; shaded = $p < 0.05$. Price-adjusted pay has no bootstrap, so that column is clustered. Tentative: $p \geq 0.02$ on the bootstrap or the remote check, or fails a check. Nominal only: fails price adjustment alone.

Three of seven pre-registered hypotheses are supported

	Hypothesis	Predicted	Estimate	p (bootstrap)	Verdict
H1	Seniority dominates advertised pay	+	+11.2% per step	<0.001	Supported
H2	A mandate raises disclosure	+	91.3% vs. 47.5%	descriptive	Supported, descriptively
H3	Required experience raises pay	+	+2.5% per year	<0.001	Supported
H4	AI/ML roles carry a premium	+	+8.0%	0.083	Inconclusive
H5	A required degree raises pay	+	−5.0%	0.074	Inconclusive
H6	Mandate states advertise wider ranges	+	−5.8% range width	0.497 (clustered)	Inconclusive
H7	Data centers pay more than utilities	+	+10.3%	0.584	Inconclusive

H5 has the wrong sign (−5.0%); clustered errors would call it significant ($p = 0.055$), the bootstrap does not ($p = 0.074$).

Early-career postings: 80 observations are too few to pin down premiums



80 early-career postings that state pay, from 36 employers

Full sample: N = 543

Crimson: full sample (core model).
Grey, hollow: early-career subsample (no bootstrap; clustered p).

At $p < 0.05$ (clustered): a stated cloud skill (+24%) and an AI/ML role family (+25%). Intervals this wide cannot be read as findings.

In the full sample, requires a degree, ML/AI skill, AI/ML role family, data-center operator and pay-transparency mandate are not significant under the bootstrap.

Early career: an entry-level title, or a stated experience minimum of 3 years or less. Bars: 95% cluster-robust CI. Seniority is left off the plot: the subsample spans only entry and mid ranks.

What did not hold up

1 Overturned by the bootstrap

ML/AI skill +9.8%: $p = 0.033$ clustered, 0.074 bootstrap
AI/ML role family +8.0%: $p = 0.048$ clustered, 0.083 bootstrap

2 A degree requirement: wrong sign, inconclusive

Predicted to raise pay; estimated -5.0% (bootstrap $p = 0.074$).
Not distinguishable from zero at this cluster count.

3 STEM degree: tentative

+5.1% (bootstrap $p = 0.043$), but fails without Crusoe ($p = 0.086$) and without federal rows ($p = 0.051$).

4 Regional premiums: nominal only

West +13.5% nominal; price-adjusted $p = 0.149$
Northeast +8.9% nominal; price-adjusted $p = 0.706$

Reported because they were tested: a null or a reversal is a result. The AI/ML premium did not hold, so it is not claimed.

Four threats, and how each is handled

1 Advertised is not realized pay

The outcome is the midpoint of the posted range; offers can be negotiated in either direction.

2 Disclosure is selected

Where no law applies only 48% of postings state pay, so pay coefficients are conditional on disclosure. Disclosure is modelled as its own outcome.

3 Associational, not causal

One cross-section with no time variation: no difference-in-differences. Employers in mandate states differ in other ways.

4 Few, uneven clusters

100 employers, the largest 9.0% of rows. Every p-value comes from the wild cluster bootstrap; Crusoe is also dropped as a check.

Also: pay is nominal in the headline model (a price-adjusted re-estimate, N = 498, is in the robustness table); scope widened in response to the data, with every change dated in the pre-registration; four Chicago-area employers on iCIMS are absent.

Conclusion

91% vs. 48%

Disclosure follows the law

Where a law requires it, 91.3% of postings state pay; elsewhere 47.5%. Stable at 43–45 points across every cut. Associational.

+11.2%

Seniority and experience set the level

Per seniority step, and +2.5% per year of required experience, as pre-registered. A cloud skill adds +12.5% (exploratory).

Did not hold

Several expected premiums

ML/AI skill, AI/ML role family: overturned by the bootstrap. Degree requirement: wrong sign. West, Northeast: nominal pay only.

Across 543 postings from 100 employers, the sharpest regularity is not how much pay is advertised, but whether it is named at all.

AI disclosure: the collection and analysis code and this deck were produced with Claude Code, Anthropic's AI coding assistant. Study design, decisions and interpretation are the author's.

APPENDIX

Appendix

Full model tables, definitions, data quality and the employer list.

Core model (pre-specified): log advertised pay

Attribute	Coef.	Std. err.	Clustered p	Bootstrap p	% effect
Seniority (per step)	0.1059	0.0079	<0.001	<0.001	+11.2%
Required experience (per year)	0.0252	0.0043	<0.001	<0.001	+2.5%
States an experience minimum	-0.0727	0.0313	0.020	0.023	-7.0%
Requires a degree	-0.0517	0.0270	0.055	0.074	-5.0%
STEM degree	0.0496	0.0230	0.031	0.043	+5.1%
Cloud skill (AWS, Azure, GCP)	0.1181	0.0328	<0.001	<0.001	+12.5%
ML/AI skill	0.0931	0.0437	0.033	0.074	+9.8%
Remote-eligible	0.0100	0.0344	0.772	0.787	+1.0%
Advertised hourly	-0.0252	0.0407	0.535	0.644	-2.5%
Pay-transparency mandate	0.0106	0.0307	0.730	0.744	+1.1%
Northeast (vs. Midwest)	0.0851	0.0307	0.005	0.019	+8.9%
South (vs. Midwest)	0.0792	0.0426	0.063	0.154	+8.2%
West (vs. Midwest)	0.1269	0.0391	0.001	0.011	+13.5%
Data-center operator	0.0983	0.0700	0.160	0.584	+10.3%
AI/ML role family	0.0765	0.0386	0.048	0.083	+8.0%

543
observations

100
employer clusters

0.589
R²

0.122
minimum detectable effect (log points)

Crimson rows: bootstrap p < 0.05. Constant 11.334. Wild cluster bootstrap: restricted, Rademacher weights, 9,999 draws, seed 20260922. Regions relative to the Midwest.

Extended model: 27 regressors (clustered SEs, no bootstrap)

Attribute	Coef.	Std. err.	Clustered p	% effect	Attribute	Coef.	Std. err.	Clustered p	% effect
Seniority (per step)	0.0852	0.0083	<0.001	+8.9%	AI/ML role family	0.0811	0.0418	0.052	+8.4%
Required experience (per year)	0.0222	0.0033	<0.001	+2.2%	Advanced degree preferred	0.0643	0.0227	0.005	+6.6%
States an experience minimum	-0.0659	0.0259	0.011	-6.4%	Leadership named	0.0540	0.0229	0.018	+5.5%
Requires a degree	-0.0367	0.0221	0.097	-3.6%	Job level (numeric title)	-0.0888	0.0130	<0.001	-8.5%
STEM degree	0.0683	0.0220	0.002	+7.1%	In a study metro	0.0348	0.0353	0.325	+3.5%
Cloud skill (AWS, Azure, GCP)	0.0937	0.0314	0.003	+9.8%	Prior internship required	-0.1847	0.0783	0.018	-16.9%
ML/AI skill	0.0859	0.0438	0.050	+9.0%	Certification required	-0.0282	0.0354	0.426	-2.8%
Remote-eligible	0.0107	0.0296	0.718	+1.1%	Python or R	-0.0014	0.0235	0.953	-0.1%
Advertised hourly	0.1573	0.0508	0.002	+17.0%	SQL	-0.0186	0.0262	0.478	-1.8%
Pay-transparency mandate	0.0157	0.0388	0.685	+1.6%	Visualization / BI	-0.0224	0.0233	0.336	-2.2%
Northeast (vs. Midwest)	0.0604	0.0280	0.031	+6.2%	Big-data tools	-0.0456	0.0336	0.175	-4.5%
South (vs. Midwest)	0.0854	0.0450	0.057	+8.9%	Teamwork named	-0.0661	0.0212	0.002	-6.4%
West (vs. Midwest)	0.0916	0.0355	0.010	+9.6%	Communication named	-0.0262	0.0224	0.244	-2.6%
Data-center operator	0.0972	0.0597	0.103	+10.2%					

N = 543, R² = 0.669. Crimson rows: clustered p < 0.05. Estimated because the sample supports it (36.2 observations per core regressor); inference in the main slides uses the core model.

Disclosure model and the price-adjusted pay model

Disclosure (LPM)	Coef.	95% CI	p
Pay-transparency mandate	0.388	[0.21, 0.56]	<0.001
Seniority (per step)	-0.002	[-0.02, 0.02]	0.869
Remote-eligible	0.083	[-0.06, 0.23]	0.271
Data-center operator	-0.071	[-0.26, 0.12]	0.463
Northeast (vs. Midwest)	-0.131	[-0.32, 0.06]	0.170
South (vs. Midwest)	-0.328	[-0.52, -0.13]	0.001
West (vs. Midwest)	-0.054	[-0.20, 0.09]	0.461

N = 729, R² = 0.335; clustered by employer.

Pay model	Nominal p (boot.)	Real coef.	Real p
Seniority (per step)	<0.001	0.111	<0.001
Required experience (per year)	<0.001	0.024	<0.001
States an experience minimum	0.023	-0.073	0.017
Requires a degree	0.074	-0.042	0.114
STEM degree	0.043	0.049	0.036
Cloud skill (AWS, Azure, GCP)	<0.001	0.092	<0.001
ML/AI skill	0.074	0.105	0.002
Remote-eligible	0.787	0.007	0.844
Advertised hourly	0.644	-0.063	0.088
Pay-transparency mandate	0.744	-0.057	0.086
Northeast (vs. Midwest)	0.019	0.012	0.706
South (vs. Midwest)	0.154	0.040	0.306
West (vs. Midwest)	0.011	0.051	0.149
Data-center operator	0.584	0.080	0.212
AI/ML role family	0.083	0.082	0.019

Real pay divides by BEA regional price parities by state (N = 498; nationwide-remote postings have no state and drop out). Crimson rows: significant in nominal pay, not in real pay.

Who discloses pay: 543 postings that state it vs. 186 that do not

Attribute	Mean, disclosed	Mean, withheld	Difference	p
Seniority (per step)	3.403	3.419	−0.016	0.900
Required experience (per year)	2.317	2.344	−0.027	0.915
States an experience minimum	0.477	0.575	−0.098	0.021
Requires a degree	0.606	0.522	+0.084	0.047
STEM degree	0.407	0.366	+0.041	0.316
Cloud skill (AWS, Azure, GCP)	0.210	0.108	+0.102	<0.001
ML/AI skill	0.359	0.274	+0.085	0.029
Remote-eligible	0.168	0.113	+0.055	0.054
Advertised hourly	0.004	0.000	+0.004	0.158
Pay-transparency mandate	0.755	0.210	+0.545	<0.001
Northeast (vs. Midwest)	0.217	0.145	+0.072	0.022
South (vs. Midwest)	0.144	0.581	−0.437	<0.001
West (vs. Midwest)	0.401	0.081	+0.321	<0.001
Data-center operator	0.140	0.183	−0.043	0.183
AI/ML role family	0.101	0.086	+0.015	0.531
Indianapolis metro	0.006	0.043	−0.037	0.015

Why it matters

Postings that state pay differ from those that do not, most of all in location and mandate coverage. Pay coefficients describe the disclosing postings only.

Disclosure rate 74.5%. Crimson rows: difference significant at 5% (two-sample test).


How the attributes are coded

Attribute	Coding rule
Pay (dependent variable)	Midpoint of the advertised range, annualized to US dollars (hourly × 2,080); logged.
Seniority (per step)	Ordinal from the title: 1 entry, 2 mid or unlevelled, 3 senior, 4 staff/principal, 5 manager, 6 director, 7 VP and above.
Required experience	Minimum years parsed from the text; 0 when unstated, with an indicator for stating one at all.
Requires a degree / STEM degree	A bachelor's or higher is required (not “or equivalent experience”); the degree must be in a STEM field.
Cloud skill / ML/AI skill	AWS, Azure or Google Cloud named; machine learning or artificial intelligence named.
Pay-transparency mandate	Any listed work location in a state whose posting law was in force on the day the posting was seen.
Region	Census region of the work location; Midwest is the reference.
Data-center operator / AI/ML role family	Employer's industry category; role family from the title taxonomy.

Coded from posting text by word-boundary patterns in config/regressors.yaml; every coded value keeps the pattern that produced it. Full definitions: docs/codebook.md.

Advertised pay, and the sample-size conditions

Range midpoints, \$25k bins (N = 543)



Mean \$151,839, median \$144,360, SD \$50,497, range \$65,000–\$365,000. Interpretable: yes.

PRE-REGISTERED CONDITIONS (ALL PASS)

36.2

observations per regressor
need ≥ 20

100

employer clusters
need ≥ 30

9.0%

largest employer's share
need $\leq 25\%$

0.122

minimum detectable effect
need ≤ 0.25 log points

Data quality: 12 rounds of hand-auditing (rounds 1–6)

#	What was audited	What it found
1	Regressor coding	Three systematic false positives, all firing on company boilerplate rather than on anything asked of the applicant
2	role_family	6 of 53 assignments wrong (89%). Four had reached a live measurement and sat in the top eleven rows by pay
3	seniority_rank, state	21 of 141 wrong (85.1%). One defect changed the headline disclosure contrast
4	The industry umbrella itself	A multi-sector consultancy supplied 22% of the sample and three rows of it were energy work. 337 postings removed; the AI-premium finding did not survive
5	The concept role screen and dedupe	All 40 titles the new matcher admitted were read: 4 false positives caught before the rebuild. One nested-location repost found in 1,940 records and collapsed
6	Run 26: all 115 added rows, then the whole corpus	The pay parser was still halving 17 Invenergy rows and recording 9 NYISO rows at their floor. Every usable "Hitachi Energy" row belonged to a sister company. Connecticut was coded as a mandate state before its law took effect. Fixing them withdrew mandate_state and region_west, which had passed the bootstrap on the unaudited data

Every defect found is pinned by a regression test built from the real title or location string that produced it. Full log: docs/audit-log.md.

Data quality: 12 rounds of hand-auditing (rounds 7–12)

#	What was audited	What it found
7	Run 27, the first live run of the description cache	The cache reused every posting it held within its window, with no disclosure drift. A second run on the same date had overwritten the first run's files, dropping a closed posting, and an edited requisition was counted twice. Both fixed at the cause
8	Run 28 (cancelled at employer 209 of 269): all 58 added rows	GE Vernova's first collection added 18 rows and a cluster; two were off-taxonomy (packaging compliance, sales). "\$1M to \$30M" project sizes were read as \$30 an hour, the only usable row of a phantom Duke cluster. The Ashby adapter had crashed on every board since the cache commit. The first next-day cache read reused 352 descriptions with no disclosure drift
9	Run 29, the national frame expansion (133 added rows)	Crusoe's salaries were read as monthly because a "\$300 per month" commuter benefit sat above them: nine rows at their ceiling or missing. Austin Energy's postings, labelled with facility names, were rejected as having no state. Six off-taxonomy roles at new employers. Crusoe alone supplied 49 usable rows, reported with a sensitivity check
10	Run 30 (28 Sep, dispatched early at the owner's request), the first collection of batches 5-12 and USAJOBS	US towns named after foreign cities read as non-US (Eversource's Berlin, CT), and a repost differing by one full stop, both fixed between runs. Location forms never met before ("Tucson, AZ - Downtown", "Washington - Pullman", "San Francisco Bay Area") rejected every posting of a newly confirmed employer (SEL) on its first collection. Eleven off-taxonomy postings at new employers (HR, labour compliance, construction, manufacturing quality, product engineering)
11	Run 31 (29 Sep, dispatched by hand when the cron did not fire), the first collection of the federal energy agencies and batches 13-15	The dotted "Washington, D.C." read as Washington State, which put two rows in the West and cost a two-site posting its DC mandate site. A retitled Workday posting lost its posting date, a Model 2 regressor. Generated lists of three names read as one phrase in the summary, paper and deck
12	Runs 32 (the 29 Sep cron, 6.5 hours late) and 34 (30 Sep, the last collection)	A lawyer admitted on "commercial development" (Bloom Energy). A pre-registered robustness check, the pay model without federal rows, computed on every run but reported nowhere; in the final data it changes a verdict. Two generated sentences overstated: a list of survivors cut at four, and a control variable described as contradicting a prediction it never had

Every defect found is pinned by a regression test built from the real title or location string that produced it. Full log: docs/audit-log.md.

The 100 employers in the pay sample (postings stating pay)

Crusoe (49)	Vantage Data Centers (7)	American Transmission Co (2)	Canadian Solar (1)
Invenergy (47)	Guidehouse (6)	Avantus (2)	Clean Power Alliance (1)
Eversource Energy (38)	X-energy (6)	Bonneville Power Administration (2)	Cologix (1)
AEP Energy (25)	California ISO (5)	Capital Power (2)	Customized Energy Solutions (1)
Xcel Energy (20)	TerraPower (5)	Clearway Energy (2)	EdgeConneX (1)
AES Indiana (18)	Valar Atomics (5)	Equinix (2)	Emporia Energy (1)
ERCOT (18)	Verse (5)	GridCARE (2)	Energy Vault (1)
GE Vernova (18)	Alliant Energy (4)	Heron Power (2)	Energy and Environmental Economics (E3) (1)
Avangrid (16)	Charles River Associates (4)	Itron (2)	Euclid Power (1)
Scout Clean Energy (11)	New Leaf Energy (4)	Kairos Power (2)	Form Energy (1)
New York ISO (10)	Rowan Digital Infrastructure (4)	Modo Energy (2)	Industrial Electric Manufacturing (1)
Nexamp (10)	Antares (3)	NREL (2)	Leeward Renewable Energy (1)
Energy Solutions (9)	Cypress Creek Renewables (3)	NTT Global Data Centers (2)	LevelTen Energy (1)
Tapestry (9)	Flexential (3)	Plus Power (2)	Madison Energy Infrastructure (1)
Ameren Illinois (8)	GridUnity (3)	Sense (2)	NRECA National Rural Electric Cooperative Association (1)
Oklo (8)	Hanwha Renewables (3)	Solar Landscape (2)	Omnidian (1)
SEL Schweitzer Engineering Laboratories (8)	SOLV Energy (3)	Sunrun (2)	Orenda (1)
Bloom Energy (7)	STACK Infrastructure (3)	WeaveGrid (2)	Origis Energy (1)
Commonwealth Fusion Systems (7)	The Brattle Group (3)	Wunder Capital (2)	QTS Data Centers (1)
E Source (7)	Tract (3)	Adapture Renewables (1)	RMI (1)
EnergyHub (7)	Voltus (3)	Ameresco (1)	Reactivate (1)
Gridware (7)	Washington Gas (3)	Argonne National Laboratory (1)	South Jersey Industries (1)
ICF (7)	Xcimer Energy (3)	Austin Energy (1)	TRIO (1)
Radiant Industries (7)	Yes Energy (3)	Brightcore Energy (1)	VELCO (1)
The Nuclear Company (7)	1st Avenue Power (2)	CARIAN (1)	Western Area Power Administration (1)

Ordered by postings. Crusoe is the largest at 9.0%; a robustness check drops it.

References

- Arnold, D., Quach, S., & Taska, B. (2025).** The impact of pay transparency in job postings on the labor market. NBER Working Paper 34480.
- Cameron, A. C., Gelbach, J. B., & Miller, D. L. (2008).** Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics*, 90(3), 414–427.
- Cullen, Z. (2024).** Is pay transparency good? *Journal of Economic Perspectives*, 38(1), 153–180.
- Cullen, Z. B., & Pakzad-Hurson, B. (2023).** Equilibrium effects of pay transparency. *Econometrica*, 91(3), 765–802.
- U.S. Bureau of Economic Analysis.** Regional price parities by state (used to price-adjust pay).

Data, code and every generated deliverable: the study repository (pre-registration, codebook, audit log and limitations in docs/).